

# TALOS

PROTECTING YOUR NETWORK

Richard Johnson  
ToorCon San Diego 2016



---

# Go Speed Tracer

---



# Introduction

- Richard Johnson
  - Research Manager
  - Cisco Talos
- Team
  - Aleksandar Nikolic
  - Ali Rizvi-Santiago
  - Marcin Noga
  - Piotr Bania
  - Tyler Bohan
  - Yves Younan
- Special Contributor
  - Andrea Allevi
- Talos VulnDev
  - Third party vulnerability research
    - 170 bug finds in last 12 months
      - Microsoft
      - Apple
      - Oracle
      - Adobe
      - Google
      - IBM, HP, Intel
      - 7zip, libarchive, NTP
  - Security tool development
    - Fuzzers, Crash Triage
  - Mitigation development
    - FreeSentry

# Introduction

- Agenda
  - Tracing Applications
  - Guided Fuzzing
  - Binary Translation
  - Hardware Tracing
- Goals
  - Understand the attributes required for optimal guided fuzzing
  - Identify areas that can be optimized today
  - Deliver performant and reusable tracing engines

# Applications

- Software Engineering
  - Performance Monitoring
  - Unit Testing
- Malware Analysis
  - Unpacking
  - Runtime behavior
  - Sandboxing
- Mitigations
  - Shadow Stacks
  - Memory Safety checkers

# Applications

- Software Security
  - Corpus distillation
    - Minimal set of inputs to reach desired conditions
  - Guided fuzzing
    - Automated refinement / genetic mutation
  - Crash analysis
    - Crash bucketing
    - Graph slicing
    - Root cause determination
  - Interactive Debugging

# Tracing Engines

- OS Provided APIs

- Debuggers

- ptrace
    - dbgeng
    - signals

- Hook points

- Linux LTT(ng)
      - Linux perf
      - Windows Nirvana
      - Windows AppVerifier
      - Windows Shim Engine
- Check out Alex Ionescu's RECON 2015 talk

- Performance counters

- Linux perf
    - Windows PDH

# Tracing Engines

- Binary Instrumentation
  - Compiler plugins
    - gcc-gcov
    - llvm-cov
  - Binary translation
    - Valgrind
    - DynamoRIO
    - Pin
    - DynInst
    - Frida and others
    - ...



# Tracing Engines

- Native Hardware Support
  - Single Step / Breakpoint
  - Intel Branch Trace Flag
  - Intel Last Branch Record
  - Intel Branch Trace Store
  - Intel Processor Trace
  - ARM CoreSight



---

# Guided Fuzzing

---



# Evolutionary Testing


- Early work was whitebox testing
- Source code allowed graph analysis prior to testing
- Fitness based on distance from defined target
- Complex fitness landscape
  - Difficult to define properties that will get from A to B
- Applications were not security specific
  - Safety critical system DoS

# Guided Fuzzing

- Incrementally better mutational dumb fuzzing
- Trace while fuzzing and provide feedback signal
- Evolutionary algorithms
  - Assess fitness of current input
  - Manage a pool of possible inputs
- Focused on security bugs

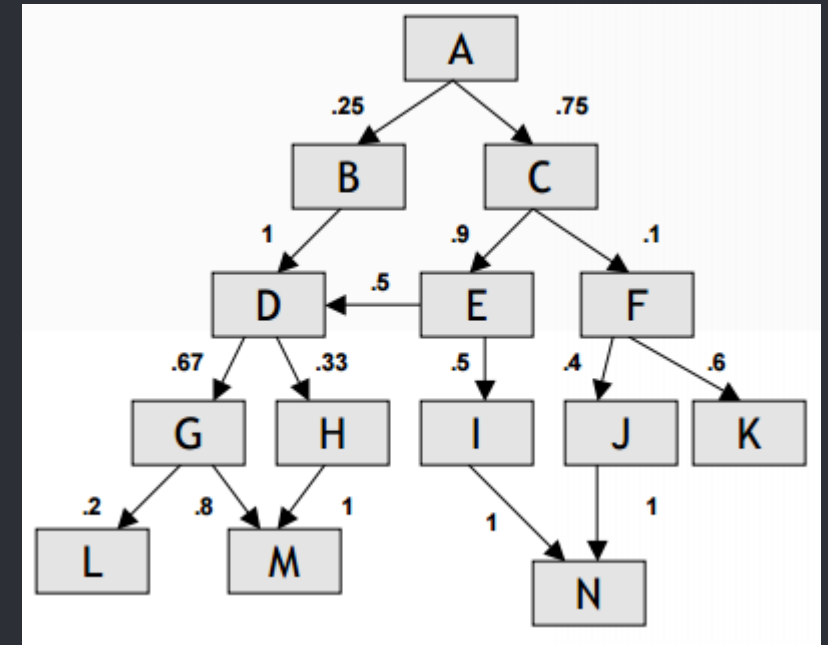
# Sidewinder

- Embleton, Sparks, Cunningham 2006
- Features
  - Simple genetic algorithm approach
    - crossover, mutation, fitness
  - Mutated context free grammar instead of sample fuzzing
  - Markov process for fitness
    - Analyzes probability of path taken by sample
  - Block coverage via debugger API
    - Reduced overhead by focusing on subgraphs



# Sidewinder

- Embleton, Sparks, Cunningham 2006
- Contributions
  - Genetic algorithms for fuzzing
  - Markov process for fitness
  - System allows selection of target code locations
- Observations
  - Never opensourced
  - Interesting concepts not duplicated



# Evolutionary Fuzzing System

- Jared DeMott 2007
- Features
  - Block coverage via Process Stalker
    - Windows Debug API
    - Intel BTF
  - Stored trace results in SQL database
    - Lots of variables required structured storage
  - Traditional genetic programming techniques
    - Code coverage + diversity for fitness
    - Sessions
    - Pools
    - Crossover
    - Mutation

# Evolutionary Fuzzing System

- Jared DeMott 2007
- Contributions
  - First opensource implementation of guided fuzzing
  - Evaluated function vs block tracing
    - For large programs found function tracing was equally effective
    - Likely an artifact of doing text based protocols
- Observations
  - Academic
    - Approach was too closely tied to traditional genetic algorithms
    - Not enough attention to performance or real world targets
    - Only targeted text protocols



# Americian Fuzzy Lop

- Michal Zalewski 2013
  - Bunny The Fuzzer 2007
- Features
  - Block coverage via compile time instrumentation
  - Simplified approach to genetic algorithm
    - Edge transitions are encoded as tuple and tracked in global map
    - Includes coverage and frequency
  - Uses variety of traditional mutation fuzzing strategies
  - Dictionaries of tokens/constants
  - First practical high performance guided fuzzer
  - Helper tools for minimizing test cases and corpus
  - Attempts to be idiot proof

# Americian Fuzzy Lop

- Michal Zalewski 2013
  - Bunny The Fuzzer 2007
- Contributions
  - Tracks edge transitions
    - Not just block entry
  - Global coverage map
    - Generation tracking
  - Fork server
    - Reduce fuzz target initialization
  - Persistent mode fuzzing
  - Builds corpus of unique inputs reusable in other workflows

american fuzzy lop 0.47b (readpng)			
<b>process timing</b>		<b>overall results</b>	
run time : 0 days, 0 hrs, 4 min, 43 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 26 sec		total paths : 195	
last uniq crash : none seen yet		uniq crashes : 0	
last uniq hang : 0 days, 0 hrs, 1 min, 51 sec		uniq hangs : 1	
<b>cycle progress</b>		<b>map coverage</b>	
now processing : 38 (19.49%)		map density : 1217 (7.43%)	
paths timed out : 0 (0.00%)		count coverage : 2.55 bits/tuple	
<b>stage progress</b>		<b>findings in depth</b>	
now trying : interest 32/8		favored paths : 128 (65.64%)	
stage execs : 0/9990 (0.00%)		new edges on : 85 (43.59%)	
total execs : 654k		total crashes : 0 (0 unique)	
exec speed : 2306/sec		total hangs : 1 (1 unique)	
<b>fuzzing strategy yields</b>		<b>path geometry</b>	
bit flips : 88/14.4k, 6/14.4k, 6/14.4k		levels : 3	
byte flips : 0/1804, 0/1786, 1/1750		pending : 178	
arithmetics : 31/126k, 3/45.6k, 1/17.8k		pend fav : 114	
known ints : 1/15.8k, 4/65.8k, 6/78.2k		imported : 0	
havoc : 34/254k, 0/0		variable : 0	
trim : 2876 8/931 (61.45% gain)		latent : 0	

# Americian Fuzzy Lop

- Michal Zalewski 2013
  - Bunny The Fuzzer 2007
- Observations
  - KISS works when applied to guided fuzzing
  - Performance top level priority in design
    - Source instrumentation can't be beat
    - Evolutionary system hard to beat without greatly increasing complexity / cost
  - Simple to use, finds tons of bugs
  - Fostered a user community
    - Developer contributions somewhat difficult
  - Current state of the art due to good engineering and feature set
  - Only mutational fuzzer system to have many third-party contributions
    - Binary support via QEMU and Dyninst
    - More robust compiler instrumentations, ASAN support

# honggfuzz

- Robert Swiecki 2010
  - Guided fuzzing added in 2015
- Features
  - Block coverage
    - Hardware performance counters
    - ASanCoverage
  - Bloom filter for trace recording
  - User-supplied mutation functions
  - Linux, FreeBSD, OSX, Cygwin support
- Contributions
  - First guided fuzzer to focus on hardware tracing support
- Observations
  - Naive seed selection for most algorithms, only the elite survive (OTTES)
    - Some modes use bloom filter
  - Easy to extend, active development

# Choronzon

- Features
  - Brings back specific genetic programming concepts
  - Contains strategies for dealing with high level input structure
    - Chunk based
    - Hierarchical
    - Containers
  - Format aware serialization functionality
  - Uses DBI engines for block coverage (PIN / DynamoRIO)
  - Attempts to be cross-platform
- Contributions
  - Reintroduction of more complex genetic algorithms
  - Robust handling of complex inputs through user supplied serialization routines
- Observations
  - Performance not a focus

# Honorable mentions

- autodafe
  - Martin Vuagnoux 2004
  - First generation guided fuzzer using pattern matching via API hooks
- Blind Code Coverage Fuzzer
  - Joxean Koret 2014
  - Uses off-the-shelf components to assemble a guided fuzzer
    - radamsa, zzuf, custom mutators
    - drcov, COSEINC RunTracer for coverage
- covFuzz
  - Atte Kettunen 2015
  - Simple node.js server for guided fuzzing
  - custom fuzzers, ASanCoverage

# Guided Fuzzing

- Required
  - Fast tracing engine
    - Block based granularity
  - Fast logging
    - Memory resident coverage map
  - Fast evolutionary algorithm
    - Minimum of global population map, pool diversity
- Desired
  - Portable
  - Easy to use
  - Helper tools
  - Grammar detection
- AFL and Honggfuzz still most practical options



---

# Binary Translation

---





# Binary Translation

- Binary translation is a robust program modification technique
  - JIT for hardware ISAs
- General overview is straightforward
  - Copy code to cache for translation
  - Insert instructions to modify original binary
  - Link blocks into traces
- Performance comes from smart trace creation
  - Originally profiling locations for hot trace
  - Early optimizations in Dynamo from HP
    - Next Executing Tail
    - Traces begin at backedge or other trace exit
  - Ongoing optimization work happens here
    - VMware - Early Exit guided

# Binary Translation

- Advantages
  - Supported on most mainstream OS/archs
  - Can be faster than hardware tracing
  - Can easily be targeted at certain parts of code
  - Can be tuned for specific applications
- Disadvantages
  - Performance overhead
    - Introduces additional context switch
  - ISA compatibility not guaranteed
  - Not always robust against detection or escape

# Valgrind

- Obligatory slide
- Lots of deep inspection tools
- VEX IR is well suited for security applications
- Slow and Linux only, DynamoRIO good replacement
- Many cool tools already exist
  - Flayer
  - Memgrind

# Pin

- “DBT with training wheels”
- Features
  - Trace granularity instrumentation
    - Begin at branch targets, end at indirect branch
  - Block/instruction level hooking supported
  - Higher level C++ API w/ helper routines
  - Closed source
- Observations
  - Delaying instrumentation until trace generation is slower
  - Seems most popular with casual adventurers
  - Limited inlining support
  - Less tuning options
  - Cannot observe blocks added to cache so ‘hit trace’ not possible

# Pin

- Example

```
VOID Trace(TRACE trace, VOID *v)
{
    for (BBL bbl = TRACE_BblHead(trace); BBL_Valid(bbl); bbl
        = BBL_Next(bbl))
    {
        BBL_InsertCall(bbl, IPOINT_ANYWHERE, AFUNPTR(basic_block_hook),
            IARG_FAST_ANALYSIS_CALL, IARG_END);
    }
}
```

# DynamoRIO

- “A connoisseur's DBT”
- Features
  - Block level instrumentation
    - Blocks are directly copied into code cache
  - Direct modification of IL possible
  - Portable
    - Linux, Windows, Android
    - x86/x64, ARM
  - C API / BSD Licensed (since 2009)
- Observations
  - Much more flexible for block level instrumentation
  - Performance is a priority, Windows is a priority
  - Powerful tools like Dr Memory
    - Shadow memory, taint tracking
    - Twice as fast as Valgrind memcheck

# DynamoRIO

- Example

```
event_basic_block(void *dcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating)
{
    instr_t *instr, *first = instrlist_first(bb);
    uint flags;
    /* Our inc can go anywhere, so find a spot where flags are dead. */
    for (instr = first; instr != NULL; instr = instr_get_next(instr))
    {
        flags = instr_get_arith_flags(instr);
        /* OP_inc doesn't write CF but not worth distinguishing */
        if (TESTALL(EFLAGS_WRITE_6, flags) && !TESTANY(EFLAGS_READ_6,
            flags))
            break;
    }
}
```

...

# DynamoRIO

- Example

```
if (instr == NULL)
    dr_save_arith_flags(drcontext, bb, first, SPILL_SLOT_1);

instrlist_meta_preinsert(bb,
    (instr == NULL) ? first : instr,
    INSTR_CREATE_inc(drcontext,
        OPND_CREATE_ABSMEM((byte *)&global_count, OPSZ_4)));

if (instr == NULL)
    dr_restore_arith_flags(drcontext, bb, first, SPILL_SLOT_1);
return DR_EMIT_DEFAULT;
}
```



# DynInst

- “Static rewriting IS possible!”
- Features
  - Static rewriting support
    - Dynamically linked binaries only
    - Eliminates issues with instruction cache misses common to DBT engines
  - Function level analysis
    - Tools must manually walk Dyninst provided CFG to instrument blocks
  - Modular C++ API / LGPL
- Observations
  - Fastest binary instrumentation out there
  - Development is slow
    - Patches we sent in for PE relocation support still not merged
  - Building Dyninst is NP-Hard
    - Use my Dockerfile on [github.com/talos-vulndev/afl-dyninst](https://github.com/talos-vulndev/afl-dyninst)

# DynInst

- Example

```
bool insertBBCallback(BPatch_binaryEdit * appBin, BPatch_function * curFunc,
                     char *funcName, BPatch_function * instBBIncFunc, int *bbIndex)
{
    unsigned short randID;
    BPatch_flowGraph *appCFG = curFunc->getCFG ();
    BPatch_Set<BPatch_basicBlock*> allBlocks;
    BPatch_Set<BPatch_basicBlock*>::iterator iter;
    for (iter = allBlocks.begin (); iter != allBlocks.end (); iter++)
    {
        unsigned long address = (*iter)->getStartAddress ();

        randID = rand() % USHRT_MAX;
        BPatch_Vector<BPatch_snippet*> instArgs;
        BPatch_constExpr bbId (randID);
        instArgs.push_back (&bbId);
    }
}
```

...

# DynInst

- Example

```
...
    BPatch_point *bbEntry = (*iter)->findEntryPoint();
    BPatch_funcCallExpr instIncExpr (*instBBIncFunc, instArgs);
    BPatchSnippetHandle *handle =
        appBin->insertSnippet(instIncExpr, *bbEntry, BPatch_callBefore,
                               BPatch_lastSnippet);
    (*bbIndex)++;
}
return true;
}
```

# Tuning Binary Translation

- Only instrument indirect branches
- Delay instrumentation until input is seen
- Only instrument threads that access the data
- Move instrumentation logic to analysis routines
  - Some APIs provide IF-THEN-ELSE analysis with optimization
- Avoid trampolines
  - Be aware of code locality and instruction cache
  - Directly inline instructions, modify AST if possible
- Inject a fork server if repeatedly executing DBT
  - See our turbotrace tool



---

# Hardware Tracing

---



# CPU Event Monitoring

- Modern CPUs contain Performance Monitoring Units (PMU)
- Model Specific Registers (MSR) used for configuration
  - Requires privileged execution (kernel or better) to access
- Types
  - Event Counters
    - Polled on-demand
  - Event Sampling (non-precise)
    - Interrupts triggered when counters hit modulus value
  - Precise Event Sampling (PEBS)
    - Uses 'Debug Store'
    - Physical memory buffers
    - Interrupt when full
- Use Linux perf / pmu-tools to experiment

# Interrupt Programming

- Interrupts - low level messaging system for system devices
  - CPU Exceptions
    - GPF, SINGLE\_STEP
  - Hardware Interrupts
    - Memory mapped or IRQ based
    - All Device I/O
  - Software Interrupts
    - System calls (int 0x80)
    - Breakpoints
- OS/hypervisor drivers required to configure interrupt handlers
  - Privileged registers or interrupt vector tables

# Interrupt Programming

- Interrupt Service Routines (ISR)
  - Registered by operating systems and drivers as callbacks
- CPU checks interrupt flag (IF) register after each instruction
  - cli and sti instructions control whether IF is checked
- CPU indexes the interrupt vector table to find appropriate handler
  - Context stored / restored while servicing interrupt
- Historically Familiar Interrupts:
  - int 1 - Single Step (TF)
  - int 3 - Single opcode, specifically designed for debugging
  - int 10h - Any Demosceners?
  - int 24h - Who remembers Error Handler  
I/O Device Specific Error  
Message  
Abort, Retry, Ignore, Fail?



# Interrupt Programming

- Programmer checklist
  - Memory must not be swapped
  - Use static variables if necessary
  - Must wrap functions with assembly
    - disable interrupts
    - push all registers
    - call interrupt handler
    - pop all registers
    - iretd

# Its a Trap

- Single Stepping
  - Enabled by setting the Trap Flag
  - After each instruction, CPU checks flag and fires exception if enabled
  - Accessible from userspace
- Branch Trace Flag
  - Modifies single step behavior to trap on branch
  - Single flag in IA32\_DEBUGCTL MSR
  - Requires kernel privileges to write to MSR
  - Windows includes a mapping from DR7 to set MSR
- SS/BTF traps are sloooooooooow, not applicable for vulnerability research

# IA32\_DEBUGCTL Register

## – MSR Address 0x1d9

- LBR [0] - Enable Last Branch Record mechanism
- BTF [1] - when enabled with TF in EFLAGS does single stepping on branches
- TR [6] - enables Tracing (sending BTMs to system bus)
- BTS [7] - enables sending BTMs to memory buffer from system bus
- BTINT [8] - full buffer generates interrupt otherwise circular write
- BTS\_OFF\_OS [9] - does not count for priv. level 0
- BTS\_OFF\_USR [10] - does not count for priv. level 1,2,3
- FRZ\_LBRS\_ON\_PMI [11] - freeze LBR stack on a PMI
- FRZ\_PERFMON\_ON\_PMI [12] - disable all performance counters on a PMI
- UNCORE\_PMI\_EN [13] - uncore counter interrupt generation
- SMM\_FRZ [14] - event counters are frozen during SMM

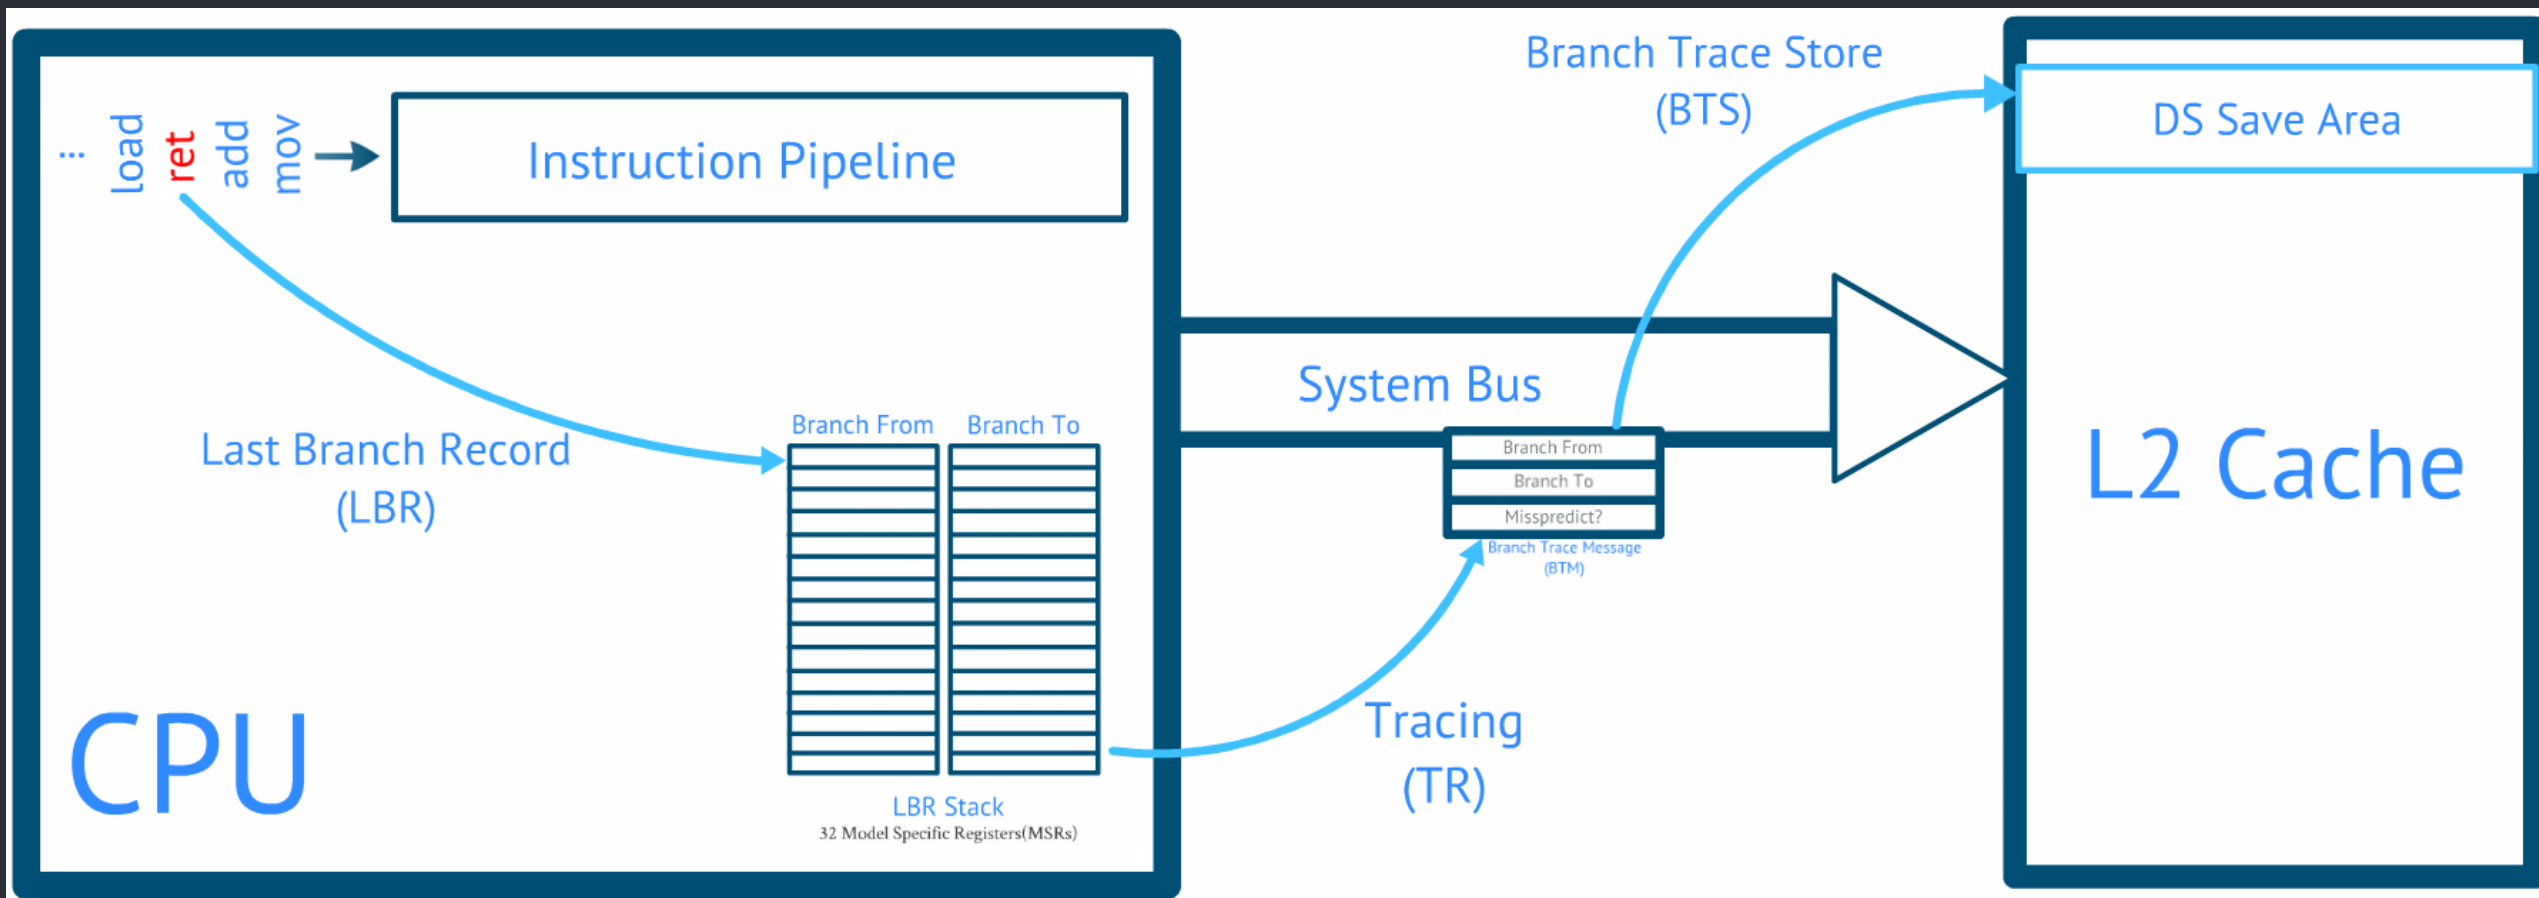
# Branch Trace Store

- First generation hardware branch tracing via PMU
- Allows configurable memory buffer for trace storage
- MSR\_IA32\_DS\_AREA MSR defines storage location
- DS\_AREA\_RECORD stored for each branch

```
struct DS_AREA {  
    u64 bts_buffer_base;  
    u64 bts_index;  
    u64 bts_absolute_maximum;  
    u64 bts_interrupt_threshold;  
    u64 pebs_buffer_base;  
    u64 pebs_index;  
    u64 pebs_absolute_maximum;  
    u64 pebs_interrupt_threshold;  
    u64 pebs_event_reset[4];  
};
```

```
struct DS_AREA_RECORD {  
    u64 flags;  
    u64 ip;  
    u64 regs[16];  
    u64 status;  
    u64 dla;  
    u64 dse;  
    u64 lat;  
};
```

# Branch Trace Store



# Branch Trace Store

- Branches in LBR registers spill to DS\_AREA
- Interrupts only when buffer is full
- Steps to enable BTS
  - Allocate memory and set MSR\_IA32\_DS\_AREA
  - Add interrupt handler to IDT
  - Register interrupt vector with APIC
    - `apic_write(APIC_LVTPC, pebs_vector);`
  - Select events with MSR\_IA32\_EVNTSEL0
    - `EVTSEL_EN | EVTSEL_USR | EVTSEL_OS`
  - Enable PEBS mode with MSR\_IA32\_PEBS\_ENABLE
  - Enable CPU perf recording with MSR\_IA32\_GLOBAL\_CTRL
- Significantly faster than BTF
- Still impractical for high speed tracing

# Intel Processor Trace

- Next generation hardware tracing support
  - Introduced in Broadwell/Skylake architecture
  - Per-hardware tracing thread
- Goal: full system branch tracing with 5-15% overhead
- Software support available in
  - Linux 4.1+ perf subsystem
  - Standalone Linux reference driver simple-pt
  - Intel VTune / System Studio\*\*
    - Remote debugging only
  - Talos IntelPT driver!

# Intel Processor Trace

- Features
  - Can trace \*SMM, HyperVisor, Kernel, Userspace [CPL -2 to 3]
  - Logs directly to physical memory
    - Bypasses CPU cache and eliminates TLB cache misses
    - Can be a contiguous segment or a set of ranges
    - Ringbuffer snapshot or interrupt mode supported
  - Minimal log format
    - One bit per conditional branch
    - Only indirect branches log dest address
    - Interrupts log source and destination
    - Decoding log requires original binaries and memory map
  - Filter logging based on CR3
  - Linux can automatically add log to coredump
  - GDB Support



# Intel Processor Trace

- 90+ pages in Intel Software Developer Manuals
- Randomly flipping bits doesn't work here ☐
- Check with CUID
- EAX = 0x14 - Intel Processor Trace
- EBX
  - Bit 00: If 1, Indicates that IA32\_RTIT\_CTL.CR3Filter can be set to 1, and that IA32\_RTIT\_CR3\_MATCH MSR can be accessed.
  - Bit 01: If 1, Indicates support of Configurable PSB and Cycle-Accurate Mode.
  - Bit 02: If 1, Indicates support of IP Filtering, TraceStop filtering, and preservation of Intel PT MSRs across warm reset.
  - Bit 03: If 1, Indicates support of MTC timing packet and suppression of COFI-based packets.
- ECX
  - Bit 00: If 1, Tracing can be enabled with IA32\_RTIT\_CTL.ToPA = 1, hence utilizing the ToPA output scheme; IA32\_RTIT\_OUTPUT\_BASE and IA32\_RTIT\_OUTPUT\_MASK\_PTRS MSRs can be accessed.
  - Bit 01: If 1, ToPA tables can hold any number of output entries, up to the maximum allowed by the MaskOffsetTableOffset field of IA32\_RTIT\_OUTPUT\_MASK\_PTRS.
  - Bit 02: If 1, Indicates support of Single-Range Output scheme.
  - Bit 03: If 1, Indicates support of output to Trace Transport subsystem.
  - Bit 31: If 1, Generated packets which contain IP payloads have LIP values, which include the CS base component
- Packet Generation (ECX = 1)
- EAX
  - Bits 2:0: Number of configurable Address Ranges for filtering.
  - Bit 31:16: Bitmap of supported MTC period encodings
- EBX
  - Bits 15-0: Bitmap of supported Cycle Threshold value encodings
  - Bit 31:16: Bitmap of supported Configurable PSB frequency encodings

# Intel Processor Trace (for programmers)

- Hardware support detection
  - CPUID with leaf 0x7 indicates support for Intel PT
  - If supported, CPUID with leaf 0x14 can return the supported PT features
- Trace Record Filtering
  - Code Privileged Level (CPL) - kernel vs userspace
  - PML4 Page Table - single process / CR3 (page-table) filtering
  - Instruction Pointer - up to 4 ranges of addresses can be specified
- Log Output Configuration
  - Single range
  - Table of Physical Addresses (ToPA)

# Intel Processor Trace (for programmers)

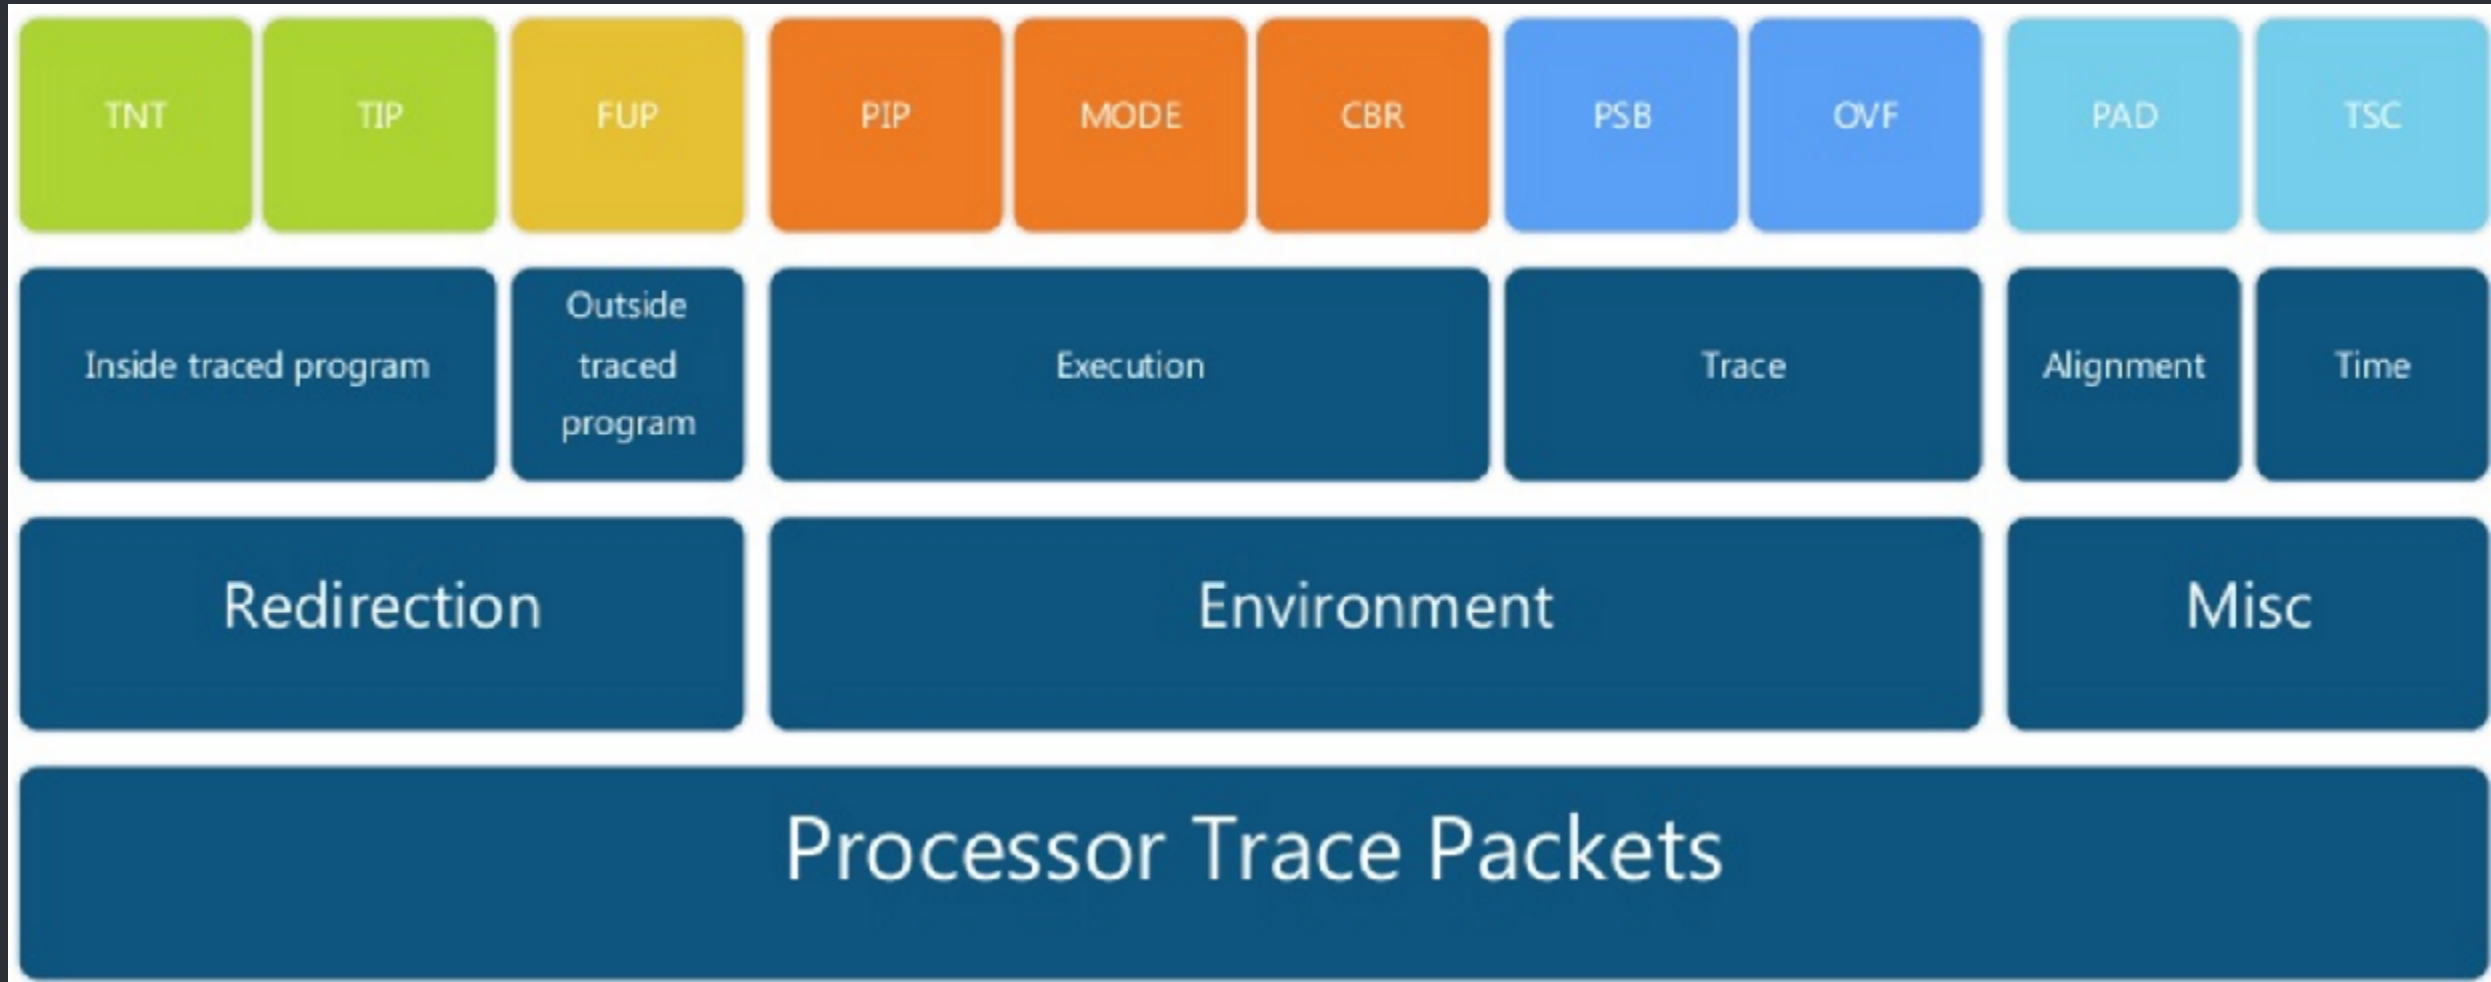
- Single Buffer Trace Logging
  - Circular or Interrupt modes (Hardware logging support)
  - Reserve memory - *MmAllocateContiguousMemory* (Windows Drivers)
  - Set the proper MSRs
    - MSR\_IA32\_RTIT\_OUTPUT\_BASE
    - MSR\_IA32\_RTIT\_OUTPUT\_MASK\_PTRS
  - Start the Tracing setting the “TraceEn” flag in the control register
  - Processor logs to in a circular-manner unless interrupt flag configured

# Intel Processor Trace (for programmers)

- Table of Physical Address (ToPA) Trace Logging
  - For large traces, non-contiguous physical memory must be used
  - ToPA is compatible with Windows Memory Descriptor List (MDL)
  - MDL is a Windows data structure for tracking physical->linear mappings

```
// Grab the physical address:  
PHYSICAL_ADDRESS physAddr = MmGetPhysicalAddress(lpBuffVa);  
// ToPA is compatible with Windows MDL data structure!  
perCpuData.u.Simple.lpTraceBuffPhysAddr = (ULONG_PTR)physAddr.QuadPart;  
  
// Allocate the relative MDL  
PMDL pPtMdl = IoAllocateMdl(lpBuffVa, (ULONG)perCpuData.qwBuffSize, FALSE, FALSE, NULL);  
if (pPtMdl) perCpuData.pTraceMdl = pPtMdl;
```

# Intel Processor Trace



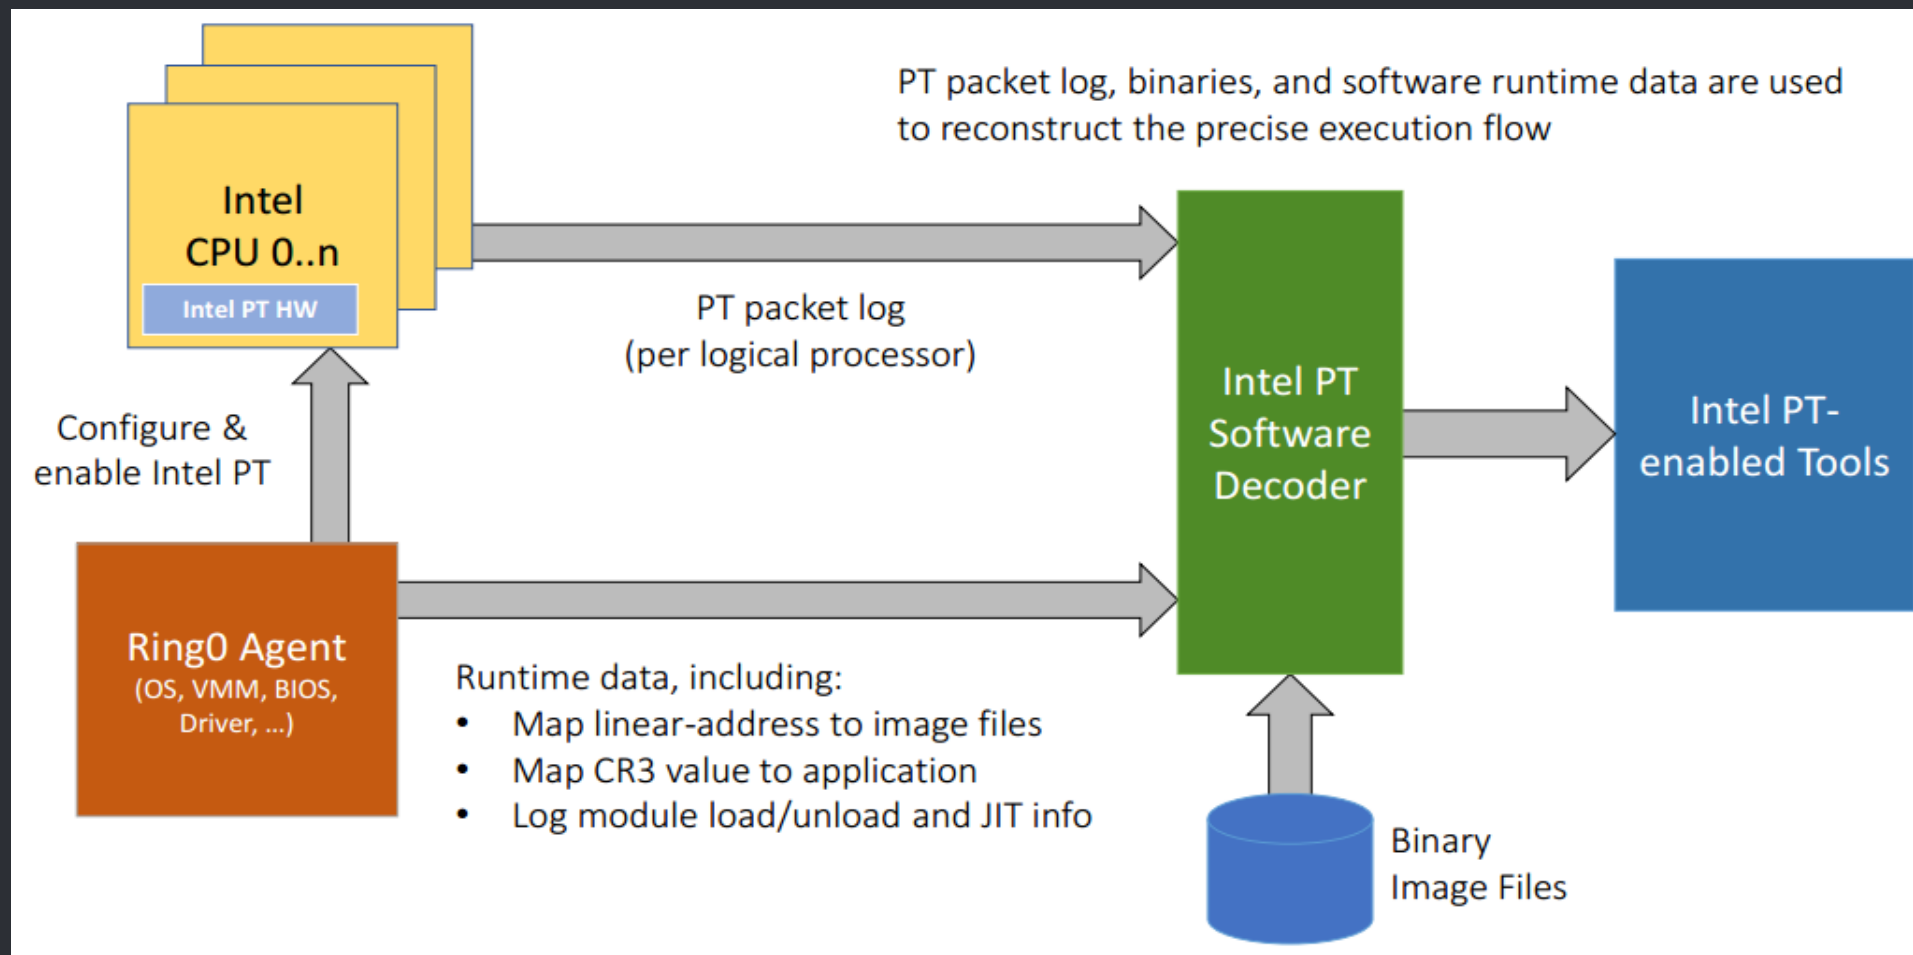
Complex log format - decode with opensource libipt library!

TALOS

# Intel Processor Trace (for programmers)

- Packet Types
  - Packet Stream Boundary (PSB)
    - Heartbeat packet generated at regular intervals (configurable)
  - Paging Information (PIP)
    - Notification of CR3 Page Table changes
  - Timing (TSC, MTC & CYC)
    - Useful for wall-clock comparisons or synchronization of logs across CPU threads
  - Control Flow (TNT, TIP, FUP)
    - TNT – Taken/Not-Taken for conditional branches
    - TIP – Taken IP address for indirect branches
    - FUP – Flow Update

# Intel Processor Trace



# Intel Processor Trace

- How to use: Linux perf tools (apt: linux-tools-common)

```
$ perf list | grep intel_pt
intel_pt//                                [Kernel PMU event]

$ perf record -e intel_pt//u date
Sun Oct 11 11:35:07 EDT 2015
[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.027 MB perf.data ]

$ perf report
...
# Samples: 1  of event 'instructions:u'
# Event count (approx.): 157207
#
# Overhead  Command  Shared Object  Symbol
# .....  .....  .....
#
# 100.00%  date      libc-2.21.so  [.] _nl_intern_locale_data
#          |
#          ---_nl_intern_locale_data
#              _nl_load_locale_from_archive
#              _nl_find_locale
#              setlocale
#
# ...
```



# Intel Processor Trace

- How to use: simple-pt reference driver

```
% sptcmd -c tcall taskset -c 0 ./tcall
cpu 0 offset 1027688, 1003 KB, writing to ptout.0
...
Wrote sideband to ptout.sideband
% sptdecode --sideband ptout.sideband --pt ptout.0 | less
TIME      DELTA  INSNS    OPERATION
frequency 32
0          [+0]    [+  1] _dl_aux_init+436
               [+  6] __libc_start_main+455 -> _dl_discover_osversio
n
...
               [+ 13] __libc_start_main+446 -> main
               [+  9]   main+22 -> f1
               [+  4]         f1+9 -> f2
               [+  2]         f1+19 -> f2
               [+  5]   main+22 -> f1
               [+  4]         f1+9 -> f2
               [+  2]         f1+19 -> f2
               [+  5]   main+22 -> f1
...

```

# Intel Processor Trace

- Talos IntelPT driver

```
struct PER_PROCESSOR_PT_DATA {  
    LPVOID lpTraceBuffVa;      // + 0x00 - VA Pointer to a contiguous memory buffer  
    ULONG_PTR lpTraceBuffPhysAddr; // + 0x08 - The physical address of the contiguous memory  
buffer  
    DWORD dwBuffSize;          // + 0x10 - The physical buffer size  
    ULONG_PTR lpTargetProcCr3; // + 0x18 - The process to monitor CR3  
};
```

# Intel Processor Trace

- Talos IntelPT driver

```
struct INTEL_PT_CAPABILITIES {  
    BOOLEAN bCr3Filtering : 1;          // [0] - CR3 Filtering Support (Indicates that  
        // IA32_RTIT_CTL.CR3Filter can be set to 1)  
    BOOLEAN bConfPsbAndCycSupported : 1; // [1] - Configurable PSB and Cycle-Accurate Mode  
    BOOLEAN bIpFiltering : 1;           // [2] - IP Filtering and TraceStop supported, and  
        // Preserve Intel PT MSRs across warm reset  
    BOOLEAN bMtcSupport : 1;            // [3] - IA32_RTIT_CTL.MTCEn can be set to 1, and MTC  
        // packets will be generated (section 36.2.5)  
    BOOLEAN bTopaOutput : 1;            // [4] - Utilize the ToPA output scheme  
    BOOLEAN bTopaMultipleEntries : 1; // [5] - ToPA tables maximum allowed (MaskOrTableOffset)
```

...

# Intel Processor Trace

- Talos IntelPT driver

```
    BOOLEAN bSingleRangeSupport : 1; // [6] - Single-Range Output Supported
    BOOLEAN bTransportOutputSupport : 1; // [7] - Output to Trace Transport Subsystem
Supported
                                                // (Setting IA32_RTIT_CTL.FabricEn to 1 is
supported)
    BOOLEAN bIpPcksAreLip : 1; // [8] - IP Payloads are LIP
    BYTE numOfAddrRanges; // + 0x01 - Number of Address Ranges
    SHORT mtcPeriodBmp; // + 0x02 - Bitmap of supported MTC Period Encodings
    SHORT cycThresholdBmp; // + 0x04 - Bitmap of supported Cycle Threshold values
    SHORT psbFreqBmp; // + 0x06 - Bitmap of supported Configurable PSB Frequency
encoding
};
```

# Intel Processor Trace

- Talos IntelPT driver

```
// Write the target CR3 value
__writemsr(MSR_IA32_RTIT_CR3_MATCH, targetCr3);

// Start tracing:
rtitCtlDesc.Fields.CR3Filter = 1;
rtitCtlDesc.Fields.FabricEn = 0;
rtitCtlDesc.Fields.Os = 0;
rtitCtlDesc.Fields.User = 1;          // Trace the user mode process
rtitCtlDesc.Fields.ToPA = 0;          // We use the single-range output scheme
rtitCtlDesc.Fields.BranchEn = 1;
//if (ptCap.bMtcSupport) {
//    rtitCtlDesc.Fields.MTCEn = 1;
//    rtitCtlDesc.Fields.MTCFreq = 10;
//}
rtitCtlDesc.Fields.TSCEn = 1;
rtitCtlDesc.Fields.TraceEn = 1;        // Switch the tracing to ON dude :-)
__writemsr(MSR_IA32_RTIT_CTL, rtitCtlDesc.All);
```

# Intel Processor Trace

- Talos IntelPT driver

```
C:\code\intelpt>instdrv.exe /I windowsptdriver.sys
C:\code\intelpt>testintelpt.exe c:\windows\system32\notepad.exe
C:\code\intelpt>..\libipt\ptdump pt_dump.bin | findstr /V pad | more
000000000000006e8 psb
000000000000006fe tsc 4e1ef46cbc
00000000000000708 cbr 1f
0000000000000070c psbend
00000000000000716 tsc 4e1ef8afb9
. . .
00000000000000ce0 cbr 1c
00000000000000cf0 tip 2: ????????4d515400
00000000000000cf5 tnt.8 ..!
00000000000000cf8 tip 2: ????????4bb10ca0
00000000000000cfd tnt.8 !!....
00000000000000cfe tnt.8 !
00000000000000d00 tip 2: ????????4d515400
00000000000000d05 tnt.8 ..!
00000000000000d08 tip 2: ????????1a91e4f0
00000000000000d0d tnt.8 !!!!!
```



---

# Outro

---



# Conclusion

- Evolutionary algorithms have a lot to offer for automation
  - <https://github.com/talos-vulndev/>
- Initial investment in development pays dividends
  - Use correct engine for long term deployment
  - Designing tracing engines is not for everyone
- Hardware tracing is approaching software performance
- This code is opensource software
  - <https://github.com/talos-vulndev/>





---

# Thank You!

---



# TALOS

[talosintel.com](https://talosintel.com)  
[blog.talosintel.com](https://blog.talosintel.com)  
[@talossecurity](https://twitter.com/talossecurity)

[@richinseattle](https://twitter.com/richinseattle)  
[rjohnson@moflow.org](mailto:rjohnson@moflow.org)

