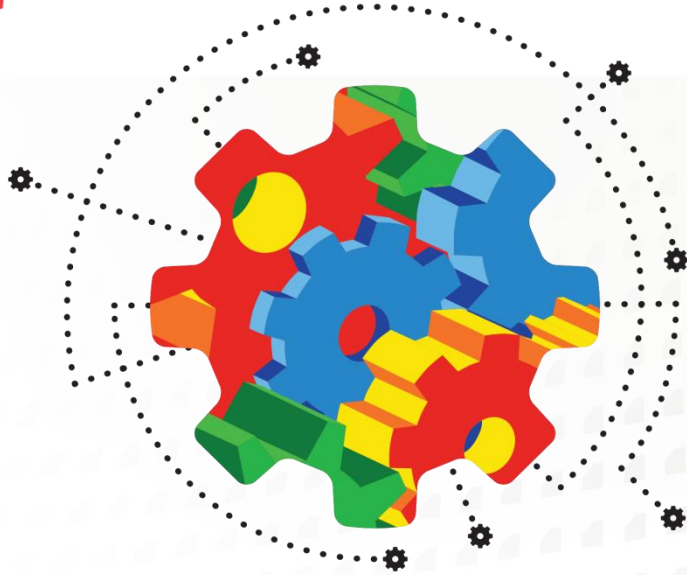


High Performance Fuzzing

Introduction

- Whoami
 - Richard Johnson / @richinseattle
 - Research Manager, Vulnerability Development
 - Cisco Talos Security Intelligence and Research Group
- Agenda
 - Why Performance Matters
 - Targeting & Input Selection
 - Engine Design
 - Host Configuration



Why Performance Matters

Why Performance Matters

- Mutational fuzzing almost seems too easy
 - Just throw some hardware at the problem
- Majority of CPU cycles are wasted
 - Program load time vs file parsing time
 - Fuzzing requires high I/O, blocking CPU
 - Mutations on large files are inefficient
- Quantitatively analyze fuzzer designs
- Qualitatively analyze fuzzer strategies

Microsoft SDL Verification Guidance

- Fuzzing is a requirement of SDLC Verification:

“Where input to file parsing code could have crossed a trust boundary, file fuzzing must be performed on that code. All issues must be fixed as described in the Security Development Lifecycle (SDL) Bug Bar. **Each file parser is required to be fuzzed** using a recommended tool.”

<https://msdn.microsoft.com/en-us/library/windows/desktop/cc307418.asp>

Microsoft SDL Verification Guidance

- Fuzzing is a requirement of SDL Verification:

“Win32/64/Mac: An **Optimized set of templates must be used**. Template optimization is based on the maximum amount of code coverage of the parser with the minimum number of templates. **Optimized templates** have been shown to **double fuzzing effectiveness** in studies. A **minimum of 500,000 iterations**, and have fuzzed at least **250,000 iterations since the last bug found/fixed** that meets the SDL Bug Bar”

<https://msdn.microsoft.com/en-us/library/windows/desktop/cc307418.asp>

Microsoft SDL Verification Guidance

- Required fuzzing is a good thing
- How did they calibrate?
 - Iterations limited by practical resources
 - Parsers with greater complexity require more resources
 - Iterations is a poor choice for defining guidance
- What properties define the theoretical limit of available resources
- What are the best practices for fuzzing to optimize our effectiveness

Historical Performance Stats

- Microsoft Windows Vista 2006
 - 15 months, 350mil iterations, 250+ file parsers
 - ~1.4mil iterations per parser (on average)
 - 300+ issues fixed
- Microsoft Office 2010
 - 800 million iterations, 400 file formats
 - 1800 bugs fixed
 - <http://blogs.technet.com/b/office2010/archive/2010/05/11/how-the-sdl-helped-improve-security-in-office-2010.aspx>
- Charlie Miller 2010
 - 3 months, 7mil iterations, 4 parsers
 - ~1.8m iterations per parser (on average)

Historical Performance Stats (cmiller)

- Charlie Miller intentionally went with a poor design
 - Microsoft minifuzz is equally stupid
- Input Selection
 - 80,000 PDFs reduced to 1515 via code coverage minset
- Targeting
 - PDF - Adobe Reader 9.2.0
 - PDF - Apple Preview (OS X 10.6.1)
 - PPT - OpenOffice Impress 3.1.1
 - PPT - Microsoft Office PowerPoint 2008 Mac

Historical Performance Stats (cmiller)

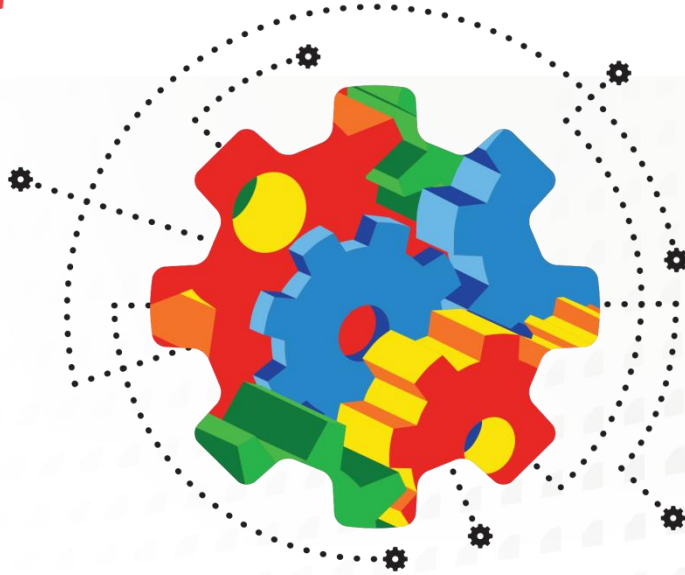
- Engine Design

- 5-lines of python to mutate input
- AppleScript to iterate files with system handler

- Performance

- About 3 weeks each target
- Reader: 3M tests, 2.2 - 8.5 sec / test
- Preview: 2.8M tests, 0.37 - 15 sec / test
- Impress: 610K tests, 4 - 60+ sec / test
- PowerPoint: 595K tests, 2 - 60 sec / test





Targeting and Input Selection

Target Selection

- 64-bit vs 32-bit applications (x86 architecture)
 - 64-bit binaries are fatter than 32-bit
 - 64-bit runtime memory usage is greater than 32-bit
 - 64-bit OSs take more memory and disk for your VMs
 - Some software only comes compiled as 32-bit binaries
 - Some fuzzers and debuggers only support 32-bit
 - 64-bit CPUs have more registers to increase performance
 - Optimization depends on compiler

Target Selection

■ So are 64-bit programs faster?

→ On x64? It varies either way to a small degree

- Chrome - Negligible
 - <http://www.7tutorials.com/google-chrome-64-bit-it-better-32-bit-version>
- Photoshop - YES?
 - 8-12% (but talks about unrelated disk i/o optimizations)
 - <https://helpx.adobe.com/photoshop/kb/64-bit-os-benefits-limitations.html>

→ On SPARC? NO

- True story, but who cares
 - http://www.osnews.com/story/5768/Are_64-bit_Binaries_Really_Slower_than_32-bit_Binaries_/page3/



Target Selection

- Much more important: Minimize lines of code
 - What is the ratio of time spent initializing program and executing actual parser code
- Optimization strategy
 - Target libraries directly
 - Write thin wrappers for each API
 - This allows feature targeting
 - Patch target to eliminate costly checksums / compression
 - This is what flayer is all about (Drewery & Ormandy WOOT'07)

Input Selection

- Input is a numerical set
- Input parsers are (should be) state machines
 - Specifications described using FSM
 - Actual parser code typically not implemented using FSM
 - LangSec Paper on high performance FSM Parsers
 - <http://www.cs.dartmouth.edu/~pete/pubs/LangSec-2014-fsm-parsers.pdf>
- Goal: search space and discover new transitions
- Each search is computationally expensive
 - We need to optimize for time



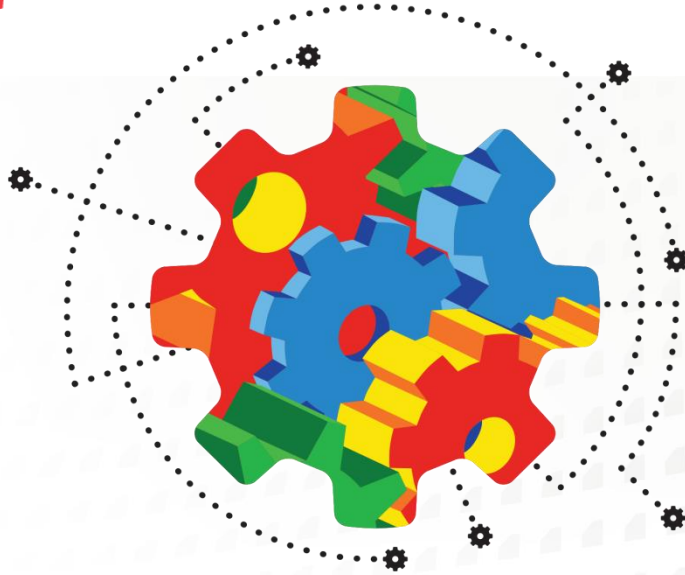
Input Selection

- Optimize input selection
 - File size is very important
 - Mutations are more meaningful with smaller input size
 - Smaller inputs are read and parsed quicker
 - Some test generation approaches utilize large amounts of memory per-input-byte
 - Specific feature set per input allows for focused targeting
 - Handcrafted or minimized samples
 - Feedback fuzzing or concolic testing automates creation of unique small inputs with different features



Input Selection

- Minset is good when it's not broken
 - Peach minset tool is not minimal set algorithm
 - Peach minset performs equivalent to random selection
- CMU Coverset
 - Optimizing Seed Selection for Fuzzing – USENIX 2014
 - <https://www.usenix.org/system/files/conference/usenixsecurity14/sec14-paper-rebert.pdf>
 - Minset helps less than expected
 - Unweighted Minset is the winner
- Charlie is redeemed, his minset tool actually works



Engine Design

Engine Design

- Generate new inputs
- Execute target with new input
- Detect failure conditions



Engine Design

- Generate new inputs
- Execute target with new input
- **Trace target execution**
- **Monitor trace output**
- Detect failure conditions
- **Detect non-failure conditions**



Input Generation

- Most important is the selection of mutators
 - AFL

Deterministic bitflip

1, 2, 4, 8, 16, 32 bits

Deterministic addition/subtraction

Values { 1 – 35 } for each byte, short, word, dword

Little endian and big endian

Deterministic 'interesting' constant values

27 boundary values

Dictionary keywords

Havoc

Random bitflips, arithmetic, block move/copy, truncate

Splice

Merge two previously generated inputs

Input Generation

- Most important is the selection of mutators
 - Radamsa

ab: enhance silly issues in ASCII string data handling
bd: drop a byte
bf: flip one bit
bi: insert a random byte
br: repeat a byte
bp: permute some bytes
bei: increment a byte by one
bed: decrement a byte by one
ber: swap a byte with a random one
sr: repeat a sequence of bytes
sd: delete a sequence of bytes
ld: delete a line

Input Generation

- Most important is the selection of mutators
 - Radamsa

lds: delete many lines
lr2: duplicate a line
li: copy a line closeby
lr: repeat a line
ls: swap two lines
lp: swap order of lines
lis: insert a line from elsewhere
lrs: replace a line with one from elsewhere
td: delete a node
tr2: duplicate a node
ts1: swap one node with another one
ts2: swap two nodes pairwise



Input Generation

- Most important is the selection of mutators

→ Radamsa

tr: repeat a path of the parse tree
uw: try to make a code point too wide
ui: insert funny unicode
num: try to modify a textual number
xp: try to parse XML and mutate it
ft: jump to a similar position in block
fn: likely clone data between similar positions
fo: fuse previously seen data elsewhere

Mutation patterns (-p)

od: Mutate once
nd: Mutate possibly many times
bu: Make several mutations closeby once



Input Generation

- Deterministic mutators first
- Permutations and infinite random mode
- Stack permutations to a reasonable level
- Need feedback loop to assess effectiveness of new mutators



Execute Target

- Using an execution loop is slow
 - process creation, linking, initialization
- Use a fork() server
 - Skip initialization
 - Copy-on-write process cloning is very fast on Linux
 - Windows and OSX manually copy process memory
 - 30x+ performance hit over COW pages



Execute Target

- Windows black magic SUA posix fork() tangent
 - ZwCreateProcess (NULL, ...) – Windows 2000
 - No sections, threads, CSRSS, User32, etc
 - RtlCloneUserProcess – Windows Vista
 - Works to limited extent
 - Applications cannot use Win32 API
 - RtlCreateProcessReflection - Windows 7
 - Designed for quick full memory dump creation
 - Does not restore threads
- Windows only does COW on mapped files by default
 - Full minidump loader an option

Execute Target

- Are you forking kidding me??

cygwin

10000 - 0m46.972s
213 exec/sec

linux

10000 - 0m2.263s
4419 exec/sec

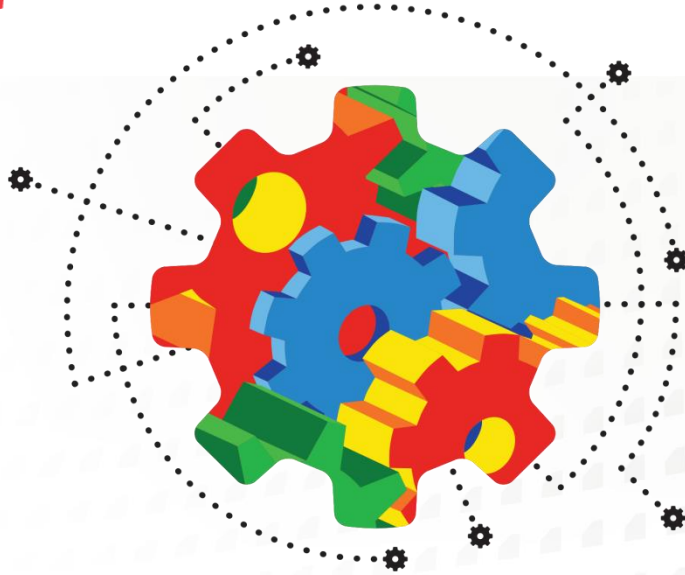
msys

10000 - 1m14.952s
135 exec/sec

native

10000 - 0m17.457s
574 exec/sec





A Forking Demo

Trace Target Execution

- Feedback loop fuzzing finally realized with AFL
 - Allows qualitative assessment of fuzzing strategy
 - Optimized instrumentation strategy
 - Optimized feedback signal
 - Source code only
- Previous attempts at binary feedback were too slow
 - EFS was overly complicated and used PaiMei
 - BCCF uses COSINC code coverage Pintool
 - Honggfuzz uses BTS



Trace Target Execution

- Hooking engine selection is critical

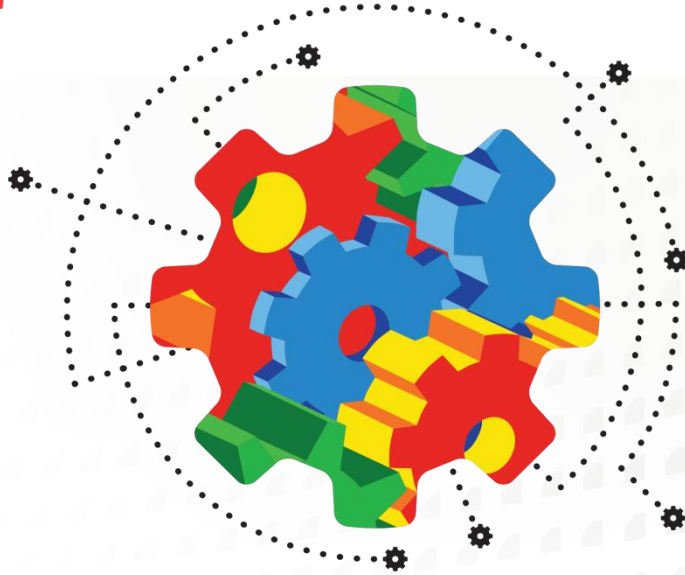
→ Pin / DynamoRIO are slow

- ~5-10x slowdown on block coverage
- Can benefit from fork server

TurboTrace:

1. Fork self in LD_PRELOADED library.
2. Ptrace the forked child.
3. Break on `_start`
4. Inject a call to the actual function that will be doing repeated `fork()`ing.
5. Step over a call.
6. Repair the `_start` and resume execution.





TurboTracer Demo

Trace Target Execution

- Hooking engine selection is critical
 - TurboTrace performance, 100 iterations
 - 20 – 50% speed increase

First test (without pintool, just instrumentation):

Pin without pintool on test_png : 55.03 seconds

Turbotrace without pintool on test_png : 37.24 seconds

Second test (bblocks pintool):

Pin bblocks pintool on test_png : 72.62 seconds

Turbotrace bblocks pintool on test_png : 51.07 seconds

Second test (calltrace pintool):

Pin calltrace pintool on test_png : 106.19 seconds

Turbotrace calltrace pintool on test_png : 85.24 seconds



Trace Target Execution

■ Hooking engine selection is critical

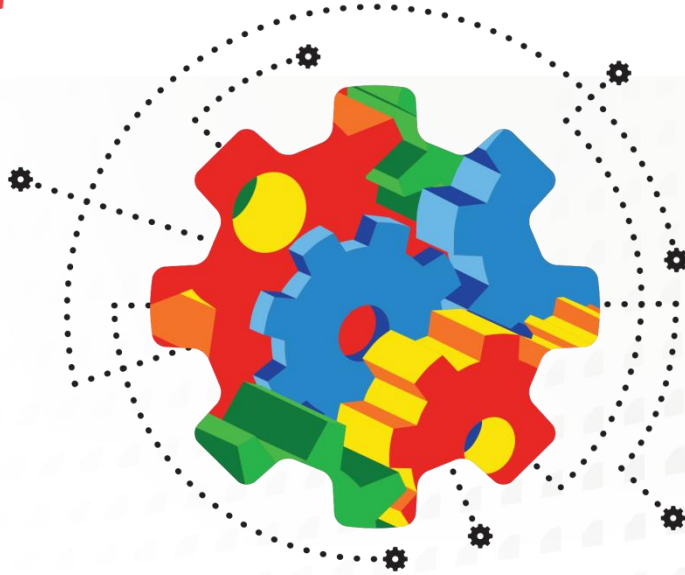
→ QEMU

- Uses QEMU userland block tracing
- Statically compiled binaries
- Linux only
- Readpng: ~860 ex/s vs ~3800 afl-gcc – 4.5x slower

→ DynInst

- Static binary rewriting
- Dynamically compiled binaries
- Linux only for now (windows port in progress)
- Readpng: ~2400 ex/s vs ~3300 afl-gcc – 1.3x slower





AFL-DYNINST DEMO

Monitor Trace Output

- Logging is critical, tracers perform way too much I/O
→ Only store enough for feedback signal
- Block coverage is weak, edge transitions are better
- Use shared memory

```
cur_location = (block_address >> 4) ^ (block_address << 8);  
shared_mem[cur_location ^ prev_location]++;  
prev_location = cur_location >> 1;
```



Detect Failure / Non-Failure

■ Failure

→ Linux

- `#define WTERMSIG(status) ((status) & 0x7f)`

→ Windows

- Debugger is the only option

■ Non-Failure

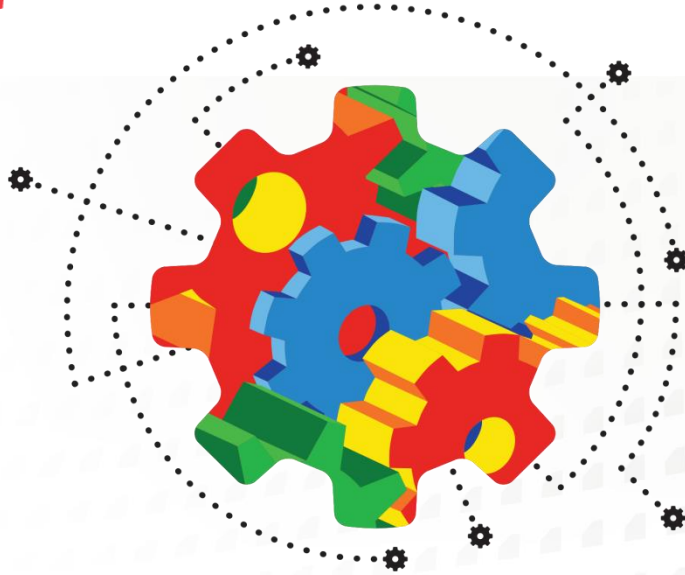
→ Timeout

- Self calibrate
- Lowest possible timeout,

→ CPU Usage

- If CPU utilization drops to near zero for X millisec





Host Configuration

System Cache

■ Windows

- Pre-Windows 7 used only 8 MB memory for filesystem cache
 - HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Control\Session Manager\Memory Management
 - Set value LargeSystemCache = 1
- Enable disk write caching in disk properties



System Cache

■ Linux

- Enables large system cache by default
- `/sbin/hdparm -W 1 /dev/hda 1` Enable write caching
- `$ sysctl -a | grep dirty`
 - `vm.dirty_background_ratio = 10`
 - `vm.dirty_background_bytes = 0`
 - `vm.dirty_ratio = 20`
 - `vm.dirty_bytes = 0`
 - `vm.dirty_writeback_centisecs = 500`
 - `vm.dirty_expire_centisecs = 3000`



Storage: HDD

- ~100 MB/s
- Cache commonly used programs proactively
 - Windows Superfetch (default)
 - Linux Preload
 - http://techthrob.com/tech/preload_files/graph.png
- Features are most useful in low memory availability scenarios
 - Typical for fuzzing w/ 1-2gb memory per VM

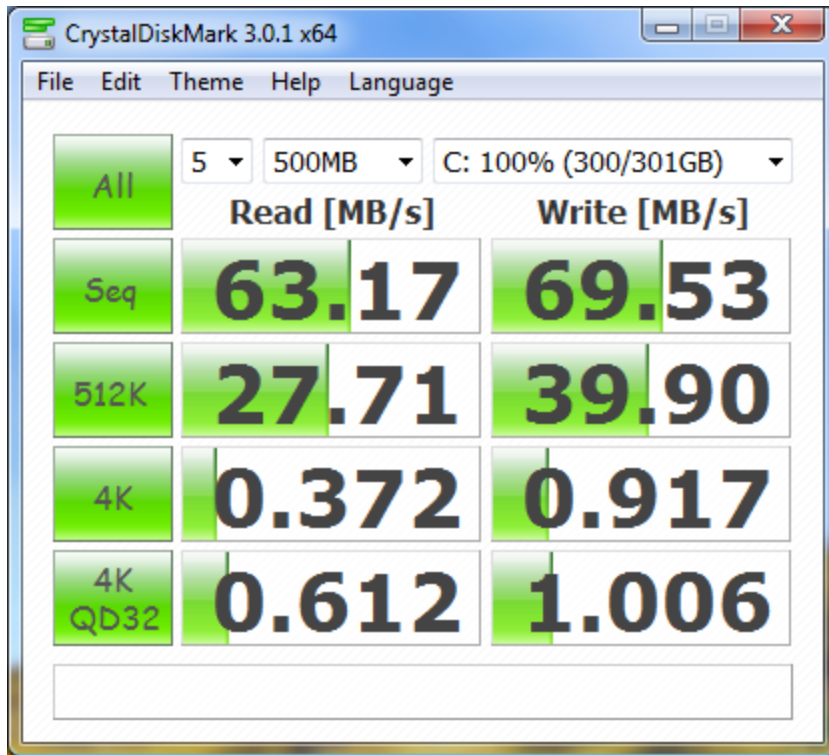


Storage: HDD

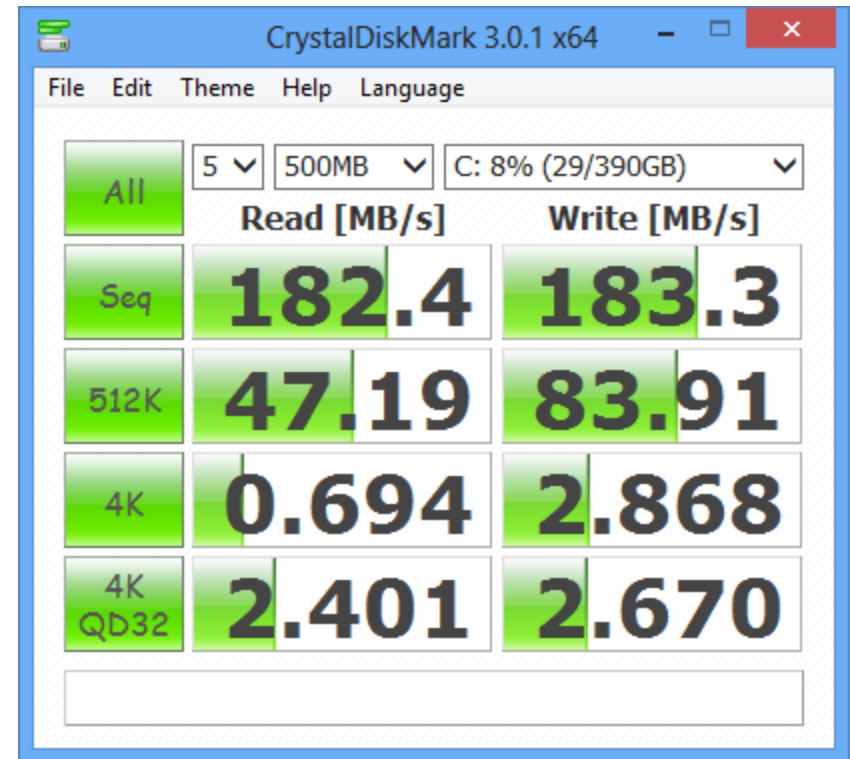
- Use a solid state USB drive for cache
 - Benefit is low latency, not high bandwidth
 - Windows ReadyBoost (available by default)
 - Random access is 10x faster on flash than hdd
 - http://www.7tutorials.com/files/img/readyboost_performance/readyboost_performance14.png
 - If you aren't already using a device for caching, and the new device is between 256MB and 32GB in size, has a transfer rate of 2.5MB/s or higher for random 4KB reads, and has a transfer rate of 1.75MB/s or higher for random 512KB write
 - <https://technet.microsoft.com/en-us/magazine/2007.03.vistakernel.aspx>
 - Linux >3.10 bache / zfs l2arc
 - 12.2K random io/sec -> 18.5K/sec with bcache, 50% increase
 - <http://bcache.evilpirate.org/>



Host Configuration



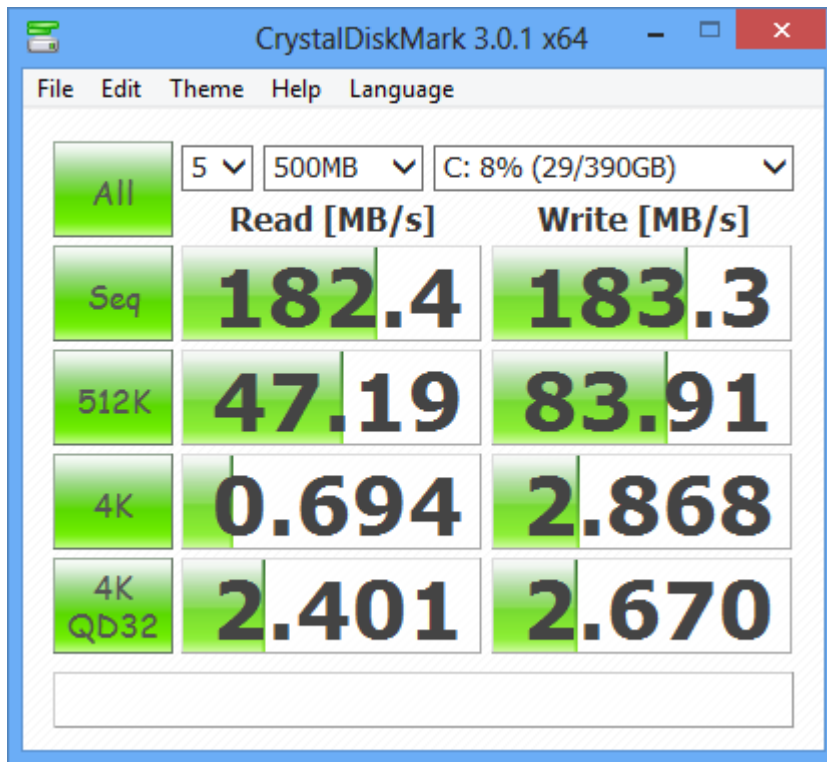
Standard HDD



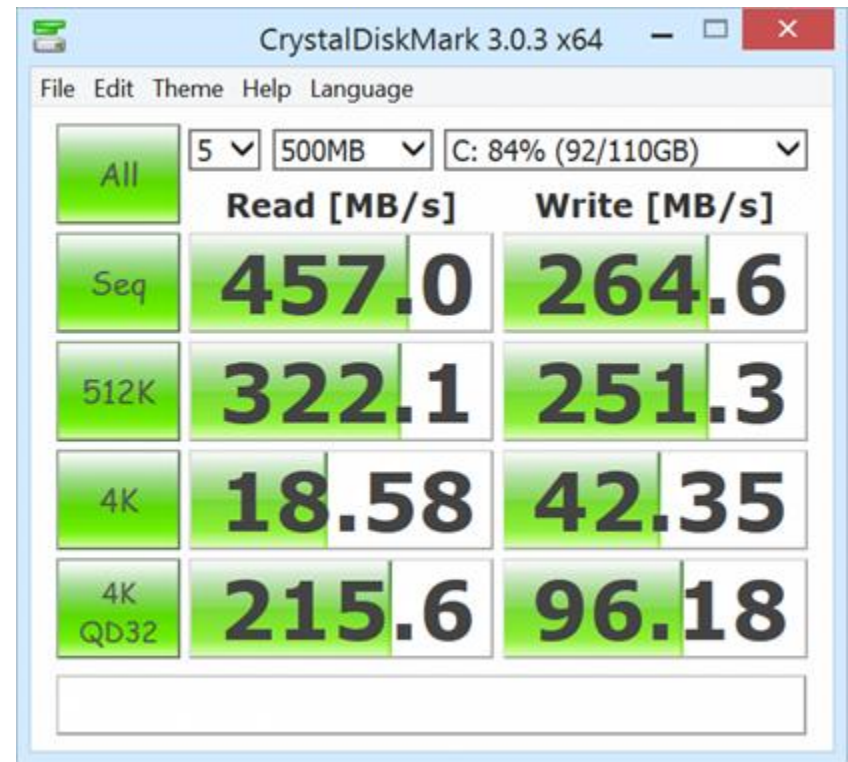
Raid 0

Storage: SSD

- Major performance gains over HDD



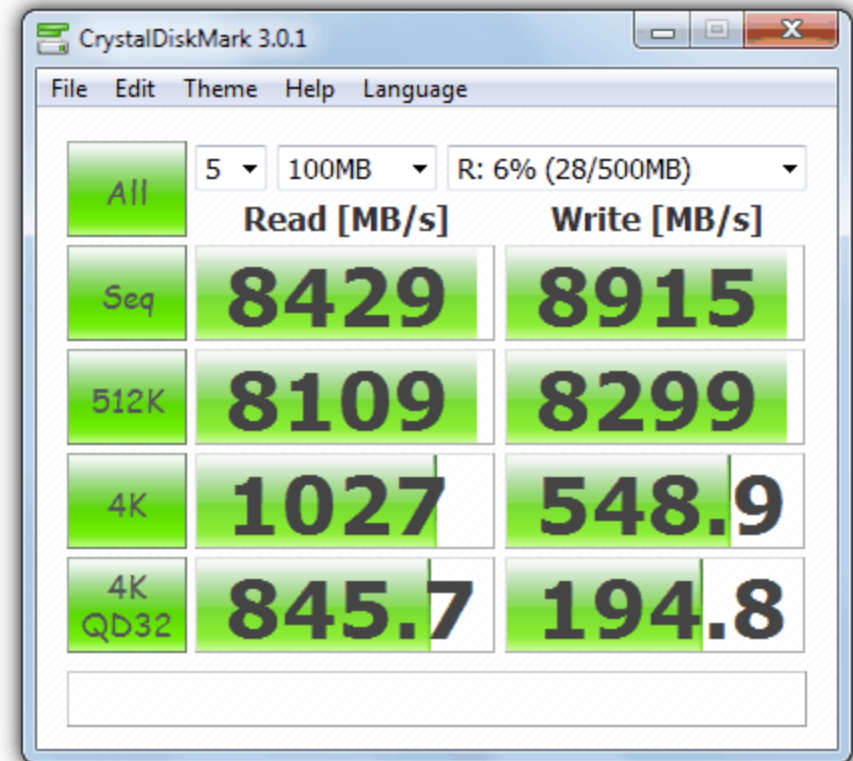
Raid 0



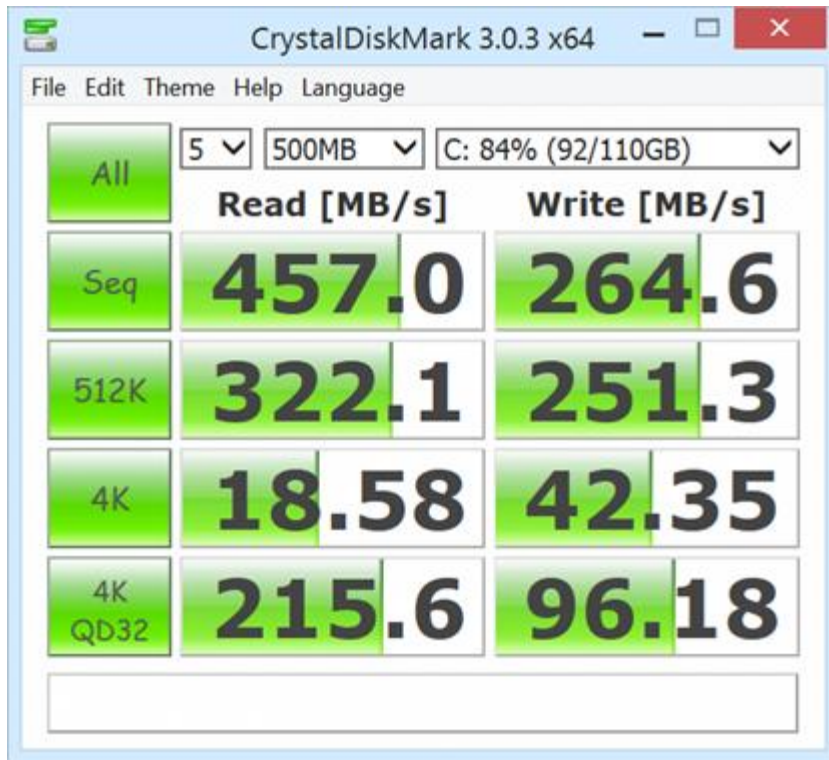
SSD

Storage: Ram Disk

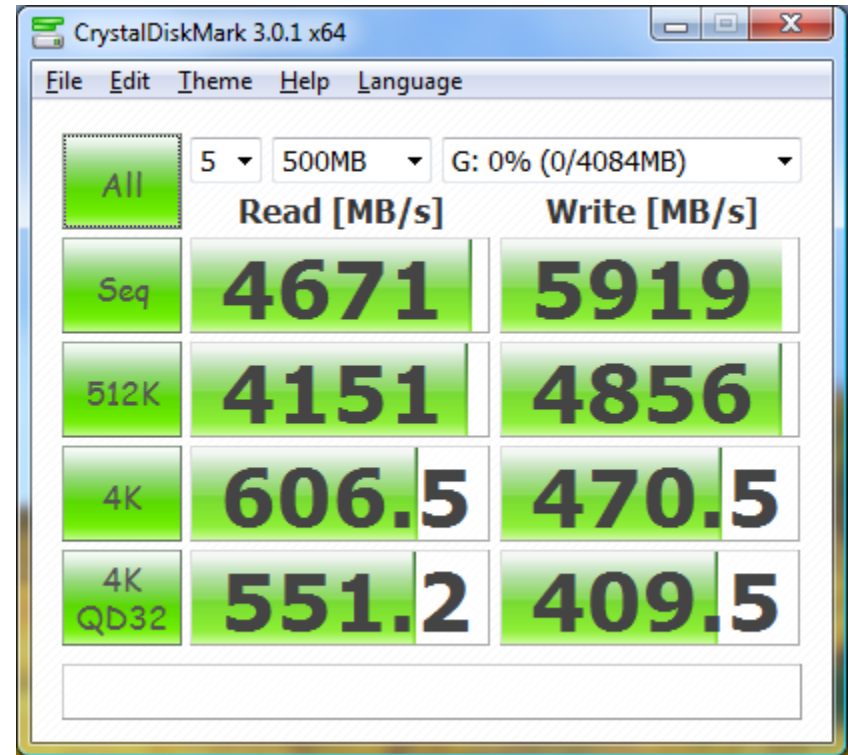
- Much faster than SSD, eliminates fragmentation
 - <http://superuser.com/questions/686378/can-ssd-raid-be-faster-than-ramdisk> (10GB/s - 17GB/s)
- Linux - built in
 - ramfs or tmpfs
- Windows - 3rd party
 - High amount of variance
 - <https://www.raymond.cc/blog/12-ram-disk-software-benchmarked-for-fastest-read-and-write-speed/>
 - SoftPerfect RamDisk is winner for free software
 - <https://www.softperfect.com/products/ramdisk/>



Host Configuration



SSD



Ramdisk

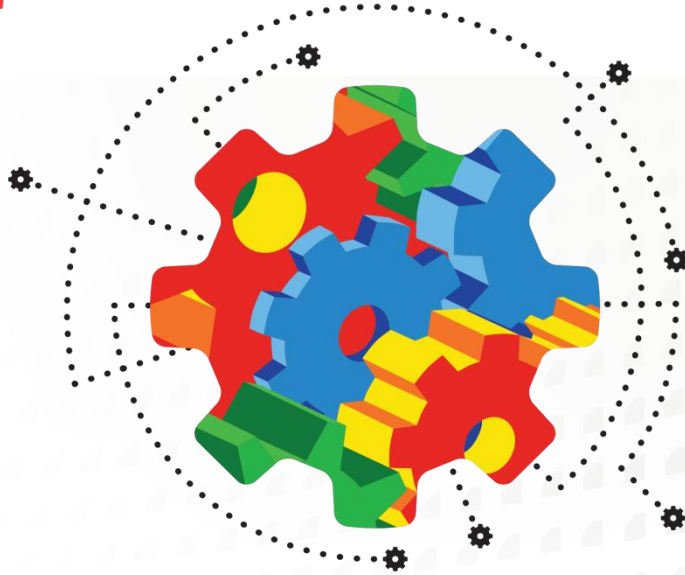
Memory

■ 32-bit memory limits

→ Linux - built in to PAE kernels

→ Windows

- Limited based on SKU of your OS
- Driver compatibility is the claimed reasoning
 - <http://blogs.technet.com/b/markrussinovich/archive/2008/07/21/3092070.aspx>
- kernel patching required
 - <http://www.geoffchappell.com/notes/windows/license/memory.htm>
 - <http://news.saferbytes.it/analisi/2012/08/x86-4gb-memory-limit-from-a-technical-perspective/>
 - <http://news.saferbytes.it/analisi/2013/02/saferbytes-x86-memory-bootkit-new-updated-build-is-out/>



Conclusions

Thank You!

- Cisco Talos VulnDev Team

- Richard Johnson

- rjohnson@sourcefire.com

- @richinseattle

- Marcin Noga

- Yves Younan

- Piotr Bania

- Aleksandar Nikolic

- Ali Rizvi-Santiago



[Concolic Testing]

- taint analysis
 - instruction level tracing
- constraint solving
 - translation to SMT
 -



Benchmark set

- AFL thoughts - found X bugs, all found within 48h
- NIST bug set -
samate.nist.gov/SARD/testcases/000/001/297/crack
[addr-bad.c](#)
- set of sliced vulns i sent before

